ECE 590.14
Midterm Report
Explorative text analysis of song lyrics using topic modelling.

Kedar Prabhudesai

October 24, 2016

1 Introduction

In this project, I have attempted to perform some explorative analysis of song lyrics data using topic modelling. Topic models use an unsupervised learning approach, i.e. they do not require manually labelled training data, and categorize commonly co-uccurring words into similar topics. Given the success of topic modelling for categorizing documents, this analysis could potentially help in classifying songs based on their lyrical content.

2 Workflow

The general workflow of the key steps involved in this project is shown in Figure 1. Broadly speaking the steps can be categorized into two main stages viz. data acquistion and data analysis.

2.1 Data acquisition

Initially I attempted to scrap lyrics data from www.azlyrics.com. I was able to extract lyrics for a few songs, but later the website stopped responding. Figuring out why I couldn't extract data from the website turned out to be challenging, and probably the most time I spent during this project. After some online research particularly on this website, I learnt that since I was making too many read requests per second, my IP address was probably blocked. Learning from my earlier mistake, I decided to write a new web scraper for extracting lyrics from www.lyricsfreak.com. This time I ensured that my read requests were spaced out, allowing around 15 seconds betwen every read request. Next, some details regarding the acquired data is provided. The following artists/bands ($n = 15$) were chosen to extract lyrics:

- Adele

- Ariana Grande

- Beatles

- Beyoncé Knowles

- Bruno Mars

- Coldplay

- Ed Sheeran

- Ellie Goulding

- John Legend

- Justin Bieber

- Justin Timberlake

- Miley Cyrus

- Rihanna

- Selena Gomez

- Taylor Swift

This is a summary of the web scraping operation. The following Python packages were used to perform these operations: `nltk, BeautifulSoup, urllib` and `re`

Step 1: The webpage associated with each artist was analysed. This webpage consisted of links to other webpages which contained song lyrics. Since the page consisted of other unwanted links too, links related to lyrics only were identified by using a simple `regexp` search operation. This way, links for first 60 song lyrics for each artist were selected.

Step 2: Using the links from the previous step, 'html' data at each link was extracted and filtered to retain only lyrical data. In this manner, lyrics from 871 songs were extracted. Since some links did not contain any lyrical information, the total number of songs was less than 900 (15×60).

Step 3: After extracting the lyrics, they we cleaned up to remove punctuations, unnecessary unicode symbols and stop words.

2.2 Data analysis using topic model

Step 1: Before running a topic model, each document i.e. lyrics for each song, was represented as bag-of-words (BoW). To get the BoW representation of each document, first the vocabulary i.e. a list of non-repeating (unique) words present in the entire corpus was determined. The total number of unique words ended up being 8700. Next, the number of instances of each word in the vocabulary per document was determined, which gives the bag-of-words representation for each document.

Step 2: I used the latent Dirichlet allocation (LDA) topic model [1] which is one of the most popular topic models used for discovering topics from text data. The LDA topic model infers the following latent structure from data:

- Topics, which are distributions over the vocabulary in the corpus.

- Per document mixing proportions of topics.

- Per word topic assignment.

I used Markov chain Monte Carlo Gibbs sampling [2] algorithm to infer the latent structure from data. Matlab implementation of the same provided by [3] was used in the analysis. For this analysis, I set the number of topics, $K = 20$. The hyperparameters for the priors in the model were chosen as a function of the number of topics.

3 Results

Figure 2 and Figure 3 show the extracted topics after running the LDA topic model. The ordering of the words is sorted in desceding order and the top 20 words in each topic are displayed. We can see from the results that the topic model does a decent job of grouping words belonging to similar topics with the following examples:

- Topic 2 seems to be related to *'love'*, with words like "love, kiss, forever, remember, never, eyes, mind, promise, lips".

- Topic 3 seems to be related to *'infatuation'*, with words like "girl, crazy, touch, whoa, damn, hot".

- Topic 6 seems to have *'artist'* related words like "bieber, beyoncé".

- Topic 7 seems to be related to *'break-up'*, with words like "can't, come, back, life, stop, please, miss, waiting, running, holding".

- Topic 11 seems to be related to *'sadness'*, with words like "dont, cry, mean, falling".

- Topic 14 seems to be related to *'partying'*, with words like "tonight, hands, babe, music, party, play, put, lover".

Having said that, there are also some topics which have words which cannot be related to each other, like the following:

- Topic 16 has words like "bad, need, good, christmas, floor, blood, wave".

- Topic 9 has words like "yeah, hard, two, die, yes, take, days".

4   Discussion

As I was analysing these results, I realized how hard a problem this is. Songs convey very complex set of emotions with commonly used everyday words. Earlier applications of topic models like scientific documents exploit the fact that specific scientific disciplines have very well defined set of words that co-occur between documents within a discipline. On the contrary, song lyrics have a lot of words which can be common between topics. For example, words like 'life', 'cry' can occur simultaneously in several topics.

5   Future Work

There are several avenues which can be explored to take this research forward. Some of which are discussed here:

- *Vary topic model parameters*: It is well known in topic modelling literature that results of topic models are very sensitive to model parameters viz. number of topics and prior hyperparameters. It would be interesting to see how topics change with these parameters.

- *Analyse lyrics in a particular genre*: I was probably too ambitious with this project to include artists and bands from a variety of genres. I would hypothesize that topics may be more closely related to songs belonging to the same genre. For example, country music may be more topics associated with 'America' and/or patriotism.

- *Augmenting non text data*: It would be interesting to see what would happen if the sound data from the actual music files can be incorporated in the analysis along with song lyrics data.
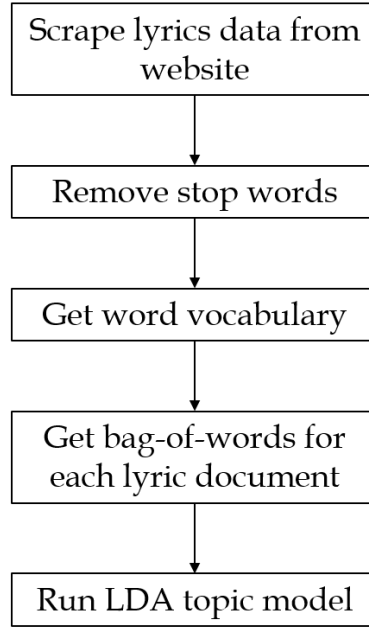
6   Figures

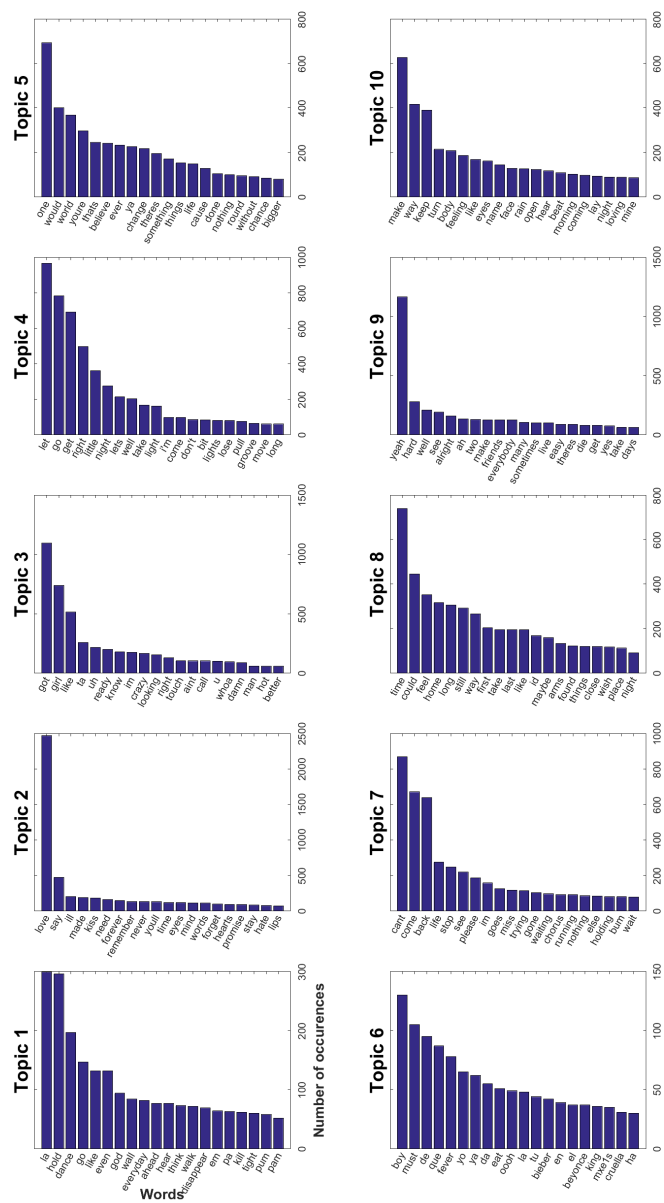Figure 1: Flowchart of various steps involved

Figure 2: Latent Dirichlet Allocation

Figure 3: Topics 1 through 10

Figure 4: Topics 11 through 20

References

[1] Blei, David M., Andrew Y. Ng, and Michael I. Jordan. "Latent dirichlet allocation." *Journal of machine Learning research* 3.Jan (2003): 993-1022.

[2] Geman, Stuart, and Donald Geman. "Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images." *IEEE Transactions on pattern analysis and machine intelligence* 6 (1984): 721-741.

[3] Griffiths, Thomas L., and Mark Steyvers. "Finding scientific topics." *Proceedings of the National academy of Sciences* 101.suppl 1 (2004): 5228-5235.